

In [1]:

```
import numpy as np
import pandas as pd
import random
import tensorflow as tf
import matplotlib.pyplot as plt
from sklearn.metrics import accuracy_score
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Flatten, Conv2D, Dense, MaxPooling2D
from tensorflow.keras.optimizers import SGD
from tensorflow.keras.utils import to_categorical
from tensorflow.keras.datasets import mnist
```

In [2]:

```
(X_train, y_train), (X_test, y_test) = mnist.load_data()
```

Type *Markdown* and LaTeX: α^2

In [3]:

```
print(X_train.shape)
```

(60000, 28, 28)

In [4]:

```
X_train[0].min(), X_train[0].max()
```

Out[4]:

(0, 255)

In [5]:

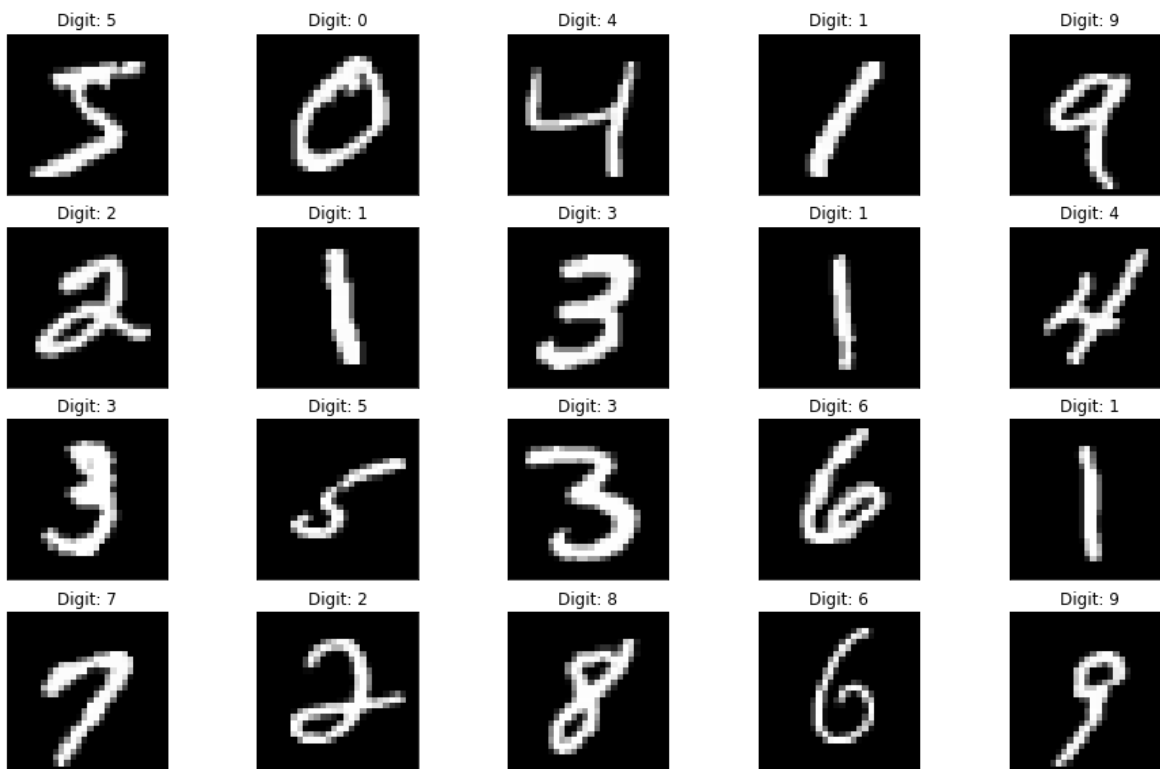
```
X_train = (X_train - 0.0) / (255.0 - 0.0)
X_test = (X_test - 0.0) / (255.0 - 0.0)
X_train[0].min(), X_train[0].max()
```

Out[5]:

(0.0, 1.0)

In [6]:

```
def plot_digit(image, digit, plt, i):
    plt.subplot(4, 5, i + 1)
    plt.imshow(image, cmap=plt.get_cmap('gray'))
    plt.title(f"Digit: {digit}")
    plt.xticks([])
    plt.yticks([])
plt.figure(figsize=(16, 10))
for i in range(20):
    plot_digit(X_train[i], y_train[i], plt, i)
plt.show()
```



In [7]:

```
X_train = X_train.reshape((X_train.shape + (1,)))
X_test = X_test.reshape((X_test.shape + (1,)))
```

In [8]:

```
y_train[0:20]
```

Out[8]:

```
array([5, 0, 4, 1, 9, 2, 1, 3, 1, 4, 3, 5, 3, 6, 1, 7, 2, 8, 6, 9],
      dtype=uint8)
```

In [9]:

```
model = Sequential([
    Conv2D(32, (3, 3), activation="relu", input_shape=(28, 28, 1)),
    MaxPooling2D((2, 2)),
    Flatten(),
    Dense(100, activation="relu"),
    Dense(10, activation="softmax")
])
```

In [10]:

```
optimizer = SGD(learning_rate=0.01, momentum=0.9)
model.compile(
    optimizer=optimizer,
    loss="sparse_categorical_crossentropy",
    metrics=["accuracy"]
)
model.summary()
```

Model: "sequential"

Layer (type)	Output Shape	Param #
=====		
conv2d (Conv2D)	(None, 26, 26, 32)	320
max_pooling2d (MaxPooling2D)	(None, 13, 13, 32)	0
flatten (Flatten)	(None, 5408)	0
dense (Dense)	(None, 100)	540900
dense_1 (Dense)	(None, 10)	1010
=====		
Total params: 542,230		
Trainable params: 542,230		
Non-trainable params: 0		
=====		

In [11]:



```
model.fit(X_train, y_train, epochs=10, batch_size=32)
```

```
Epoch 1/10
1875/1875 [=====] - 43s 22ms/step - loss: 0.2322 -
accuracy: 0.9304
Epoch 2/10
1875/1875 [=====] - 42s 22ms/step - loss: 0.0777 -
accuracy: 0.9767
Epoch 3/10
1875/1875 [=====] - 38s 20ms/step - loss: 0.0508 -
accuracy: 0.9848
Epoch 4/10
1875/1875 [=====] - 40s 21ms/step - loss: 0.0374 -
accuracy: 0.9882
Epoch 5/10
1875/1875 [=====] - 44s 23ms/step - loss: 0.0275 -
accuracy: 0.9914
Epoch 6/10
1875/1875 [=====] - 43s 23ms/step - loss: 0.0212 -
accuracy: 0.9939
Epoch 7/10
1875/1875 [=====] - 45s 24ms/step - loss: 0.0153 -
accuracy: 0.9952
Epoch 8/10
1875/1875 [=====] - 42s 23ms/step - loss: 0.0120 -
accuracy: 0.9964
Epoch 9/10
1875/1875 [=====] - 42s 22ms/step - loss: 0.0085 -
accuracy: 0.9979
Epoch 10/10
1875/1875 [=====] - 39s 21ms/step - loss: 0.0059 -
accuracy: 0.9986
```

Out[11]:


```
<keras.callbacks.History at 0x1f1cb468fa0>
```

In [12]:



```
plt.figure(figsize=(16, 10))
for i in range(20):
    image = random.choice(X_test).squeeze()
    digit = np.argmax(model.predict(image.reshape((1, 28, 28, 1))))[0], axis=-1)
    plot_digit(image, digit, plt, i)
plt.show()
```

```
1/1 [=====] - 3s 3s/step
1/1 [=====] - 0s 50ms/step
1/1 [=====] - 0s 34ms/step
1/1 [=====] - 0s 39ms/step
1/1 [=====] - 0s 32ms/step
1/1 [=====] - 0s 46ms/step
1/1 [=====] - 0s 30ms/step
1/1 [=====] - 0s 35ms/step
1/1 [=====] - 0s 31ms/step
1/1 [=====] - 0s 32ms/step
1/1 [=====] - 0s 86ms/step
1/1 [=====] - 0s 32ms/step
1/1 [=====] - 0s 33ms/step
1/1 [=====] - 0s 33ms/step
1/1 [=====] - 0s 38ms/step
1/1 [=====] - 0s 33ms/step
1/1 [=====] - 0s 32ms/step
1/1 [=====] - 0s 31ms/step
1/1 [=====] - 0s 35ms/step
1/1 [=====] - 0s 32ms/step
```



In [13]:

```
predictions = np.argmax(model.predict(X_test), axis=-1)
accuracy_score(y_test, predictions)
```


313/313 [=====] - 2s 7ms/step

Out[13]:

0.9862

In [14]:

```
n=random.randint(0,9999)
plt.imshow(X_test[n])
plt.show()
```



In [15]:

```
predicted_value=model.predict(X_test)
print("Handwritten number in the image is= %d" %np.argmax(predicted_value[n]))
```

313/313 [=====] - 2s 7ms/step
Handwritten number in the image is= 1

In [16]:



```
score = model.evaluate(X_test, y_test, verbose=0)
print('Test loss:', score[0]) #Test Loss: 0.0296396646054
print('Test accuracy:', score[1])
```

Test loss: 0.04130799323320389

Test accuracy: 0.9861999750137329

In [17]:



```
#The implemented CNN model is giving Loss=0.04130799323320389 and
#accuracy: 0.9861999750137329 for test mnist dataset
```